

Generative AI

The 5 AI Tensions Leaders Need to Navigate

by Rebecca Hinds and Robert I. Sutton

December 10, 2025




Illustration by Chen Wu

Summary. The introduction of AI into the workplace inherently creates tension. The same tools that relieve drudgery and make work easier, for example, can also

remove the challenging friction that gives work its meaning, builds crucial skills, and increases... [more](#)

[Ler em português](#)

When Polish endoscopists began using AI to detect cancer, their [accuracy improved](#). But their performance on non-AI procedures got worse. When students used AI to draft SAT-style essays, their creativity initially spiked. Yet those who started with AI-generated ideas showed [reduced alpha-wave activity](#) (a marker of creative flow), “[tended to converge on common words and ideas](#),” and their “output was very, very similar” to one another’s. And in a [2025 study](#) spanning 20 European countries, workers in highly automated jobs reported less purpose, less control, and more stress, even when their work became technically easier.

Together, these studies capture some of the many dualities of using AI at work. The same tools that relieve drudgery can also remove the challenging friction that gives work its meaning. Such tools can make work easier but sand away the struggle that builds skill and satisfaction.

We wanted to understand how these tensions play out inside real organizations—and how smart leaders are building systems and guardrails to navigate them. Amid the hype and hand-waving, we



set out to discover what's actually working. We gathered studies and lessons from more than 100 people (interviewing 35 of them) within our networks who are helping their own and other organizations bring AI into the workplace, or studying how the challenges and solutions are unfolding. These builders, executives, investors, advisors, and researchers from across the globe vary in their areas of expertise and optimism about AI in the workplace. Yet all are people we admire for their wisdom and common sense.

The result is the “[AI Transformation 100](#),” an annotated list of 100 ideas on how AI *just might* accentuate the positive and dampen the negative in our workplaces. This is our first project from the new [Work AI Institute at Glean](#), an enterprise AI platform that helps teams find knowledge, take intelligent action, and automate work. Our goals were to begin to separate experimentation from execution and hype from tangible transformation, as well as to understand where AI is driving progress, where it's getting stuck, and where it's quietly (or loudly) backfiring.

There was much noise and many inconsistencies in the studies and stories we gathered to create the AI Transformation 100. Yet one pattern stood out. The best leaders don't talk about AI in absolutes. They don't treat it as savior or saboteur. They stay curious...and a little skeptical. They ask: Are we going too far, or



not far enough? Throughout our research, five tensions kept surfacing, each an AI balancing act every leader needs to navigate.

Tension 1: Experts vs. Novices

AI is upending who gets to be an expert. What once required years of training—writing code, analyzing data, drafting legal language—can now be done by almost anyone with the right prompt. The line between expert and novice is blurring.

That shift brings both opportunity and risk. It expands who gets to contribute but also makes it easier to mistake AI fluency for mastery of a craft or subject area. The challenge isn't choosing between experts and novices; it's knowing when to lean on each.

On the one hand, experts bring rigor and depth. They've seen the failures, the edge cases, the reasons things don't work. That knowledge is essential, but it can also shut down creativity too soon. Duolingo board member John Lilly told us about two non-engineers with zero chess experience who used AI to prototype a learn-to-play [chess course](#) in just four months that outpaced other internal Duolingo initiatives. "If you bring experts in too early," he said, "they'll tell you all the reasons it won't work. AI let the non-engineers show what was possible, fast." And at Google, teams



have shifted from lengthy documents that detail product requirements to a “prototype-first” approach. With AI-powered “vibe coding,” multiple insiders tell us that they now build working demos before drafting proposals, speeding iteration and keeping good ideas from dying in committee.

But swinging too far in the other direction brings new risk. As both leaders on the front lines and researchers emphasized to us, novices using AI to imitate expertise can create “AI slop”—outputs that look convincing but collapse under scrutiny. Stanford research on which software developers are getting jobs (and which aren’t) suggests that companies are relying on experts to battle such slop. While hiring for entry-level developers has declined, demand for senior engineers continues to rise. AI can generate a passable draft, but it can’t replicate veteran developers’ judgment, elegance, or systems thinking.

To tap the imagination of novices without opening up the floodgates to AI slop, try these:

Let generalists start, but not finish. Use AI to lower the barrier to entry for early prototypes, but make sure experts step in to test, refine, and scale what works. At Stitch Fix, former chief algorithms officer Eric Colson described how algorithms scanned inventory and customer preferences to flag unmet needs (such as



styles, colors, or fabrics missing from the lineup). Custom algorithms developed by Colson's team then generated design suggestions based on those gaps. But instead of letting the system greenlight production, Stitch Fix routed those suggestions to human designers, who decided which ones would dovetail with the brand, meet quality standards, and resonate with customers. The algorithms expanded the set of creative options; the experts cut it back to what was worth doing.

Don't leave out your best people from training the model. TELUS VP of AI Alexandre Guilbault warned us about keeping high performers overly tethered to their day-to-day jobs. He said, "The best people are the ones who can drive the biggest transformation...but often organizations want to keep their best folks in operations." Leaders need to bring their top employees—such as clinicians, technicians, HR professionals, and data experts—in from the trenches to test AI models and participate in pilots from the start, even if it slows short-term execution. Leaving them out risks building systems around the habits of average performers rather than the actions that make your best people great.

Embed experts in local teams. At Glean, the "Glean on AI" team is composed of experts who represent and are formally part of different functions in the organization, such as marketing, finance, and engineering. These experts are responsible for



identifying manual processes within their specialized functions that are worth automating and creating roadmaps for turning such processes into AI-driven agents. The “AI Outcomes” team at Glean applies a similar playbook by embedding sales employees in customers’ organizations to understand how they work and develop solutions *with* them, not *for* them.

Tension 2: Centralized vs. Decentralized

As AI spreads, leaders face a familiar question: How much control should stay at the top and the center—and how much freedom belongs at the edges and lower levels of the organization?

Some organizations double down on centralization, creating AI centers of excellence to enforce standards, manage risk, and control access. But when every pilot must crawl through a maze of legal, security, and procurement reviews, innovation dies in committees, bottlenecks, and botched handoffs. By the time approvals land, the technology has already moved on.

On the other hand, when AI development becomes massively decentralized, innovation can outpace integration. Teams sprint in different directions, dashboards multiply, and automated solutions overlap and sometimes clash. The resulting overload and relentless toggling between apps and fragmented tools fuels what UC Santa Barbara professor Paul Leonardi calls digital exhaustion.



To build in enough governance without grinding every good idea to a procedural halt, try these:

Centralize for scale and safety; decentralize for learning and speed. At one U.S. university, the central IT center of excellence manages high-risk domains—data governance, integration, and infrastructure—to ensure consistency and security. Individual schools run decentralized pilots to move faster and test ideas. The CIO at the university suggested using the following rule of thumb to us:

- High risk, low experimentation (e.g., data security, compliance): Centralize.
- Low risk, high experimentation (e.g., workflow automations, prompt testing): Decentralize

Beware of shiny new roles. In moments of transformation, it's tempting to centralize progress by inventing new AI titles, such as chief AI officers, prompt engineers, and AI ethics leads. But too often, these roles become symbolic, not strategic—window dressing without budget, authority, or real influence.

Multiple leaders and academics explained to us how centralizing AI leadership under one or two overlords with new titles can slow learning and create bureaucratic choke points, whereas excessive



decentralization can scatter ownership across too many lightweight, uncoordinated, and sometimes warring roles. Instead of adding more layers, leaders, teams, and titles that add unnecessary complexity or that insiders and outsiders see as hollow symbols, distribute AI responsibilities into existing teams and roles, where authority, understanding of the work, and trust already live.

Choose technology that enforces enterprise-grade governance but enables people to do their work unfettered by friction. The best AI platforms provide guardrails (e.g., security, permissions, and audit trails) while letting individual teams safely build, test, and iterate within their own domains. Before implementing an internal AI-powered search platform, HR teams at Booking.com were “extremely nervous” about introducing AI. They worried it might expose sensitive employee information to colleagues when they used the tool to search for company resources such as HR policies and forms. As senior engineering manager Tadeu Faedrich explained, “We didn’t want people finding documents they shouldn’t have access to.” They were careful to implement a system that ensures AI searches only reveal information that employees have permission to view or access. That way, employees can freely use the tool as part of their day-to-day work without getting bogged down by approvals or worrying about privacy risks.



Tension 3: Flatter vs. Taller Hierarchy

Organization charts with too many management layers stifle momentum, trapping decisions multiple levels above where the best data, experience, and judgment live. Many organizations—especially tech companies—are jumping on the “flatter is better” bandwagon. As AI automates reporting and routine decisions, many leaders assume that they can remove layers and enable their companies to move faster. But hierarchy isn’t synonymous with stifling bureaucracy; it’s how organizations harness and make sense of complexity. Cutting layers without redesigning how work flows through the system simply shifts and often magnifies the destructive friction.

One of our interviewees, Michael Arena, former chief talent officer at General Motors and a veteran social network researcher, found that excessive flattening results in overloaded managers who create bottlenecks as they fail to keep up with decisions, crises, and day-to-day duties. Specifically, when managers lead teams with more than seven direct reports, they often work 10- to 13-hour days and still struggle to complete their work, leaving them feeling overwhelmed and burned out.

To decide whether it’s wise to flatten your org chart given the messy reality of how your work gets done, consider these:



Beware of the “flattening is faster” half-truth. Many companies have flattened their organizations to move faster, but Arena and his colleagues’ research shows how this strategy can backfire, overloading managers and slowing coordination. So, before flattening, Arena suggests assessing how much people in your organization work in two modes:

- **Heads-down work** that doesn’t require nuanced coordination and collaboration, such as coding, writing, or data analysis.
- **Heads-up work** that requires (often real-time) interdependence and two-way communication, such as building relationships with customers, brainstorming new products and services with a team, or generating and communicating the organization’s strategy.

If most of the work is heads-down execution by people who can work without tight coordination and constant communication, use AI agents to take on routine tasks and lighten the load on managers—enabling them to lead larger teams. If the work is more heads-up, keep the number of people reporting to each manager smaller, and reinvest efficiency gains from AI so managers can spend more time communicating, preventing destructive conflict, exercising judgment, and strengthening relationships.



Use AI to lighten, not eliminate, management. Offload repetitive tasks and administrative sludge—such as status updates, scheduling, and reporting—so managers can focus on what only humans do well: coaching, coordination, and judgment. At Workday, VP of people analytics Phil Wilburn told us that he stopped requiring teams to build briefing decks and weekly update docs. Instead, unstructured data from Slack conversations and project plans now flows into a single AI system. Before meetings, Wilburn asks AI to compile a brief or dig deeper into a topic. As a result, his team no longer wastes hours piecing together updates for him or answering his quick questions. AI hasn't replaced management—it's removed the administrative sludge that gets in the way of it.

Tension 4: Fast vs. Slow

The leaders of just about every company want to be seen as “doing AI”—and doing it *fast* has emerged as a sign they're doing it *right*. But faster isn't always better. Hype and smart talk can be produced far faster than the slow, painstaking work of rewiring how people actually work. The real challenge isn't moving faster and faster but knowing where to move fast and where to slow down.

In many organizations, excessive speed and impatience fuels decision-implementation gaps: Leaders move quickly to adopt



new tools but don't slow down enough to fix broken systems or understand how the technology fits within existing constraints, habits, and workflows. As a result, decisions to “do something new” take far longer to implement than expected (if they're implemented at all). Adoption becomes uneven, delayed, or quietly abandoned. AI's promise of improved efficiency and reliability fades when leaders try to jam it onto organizations that lack the systems, incentives, and skills required to put it into practice.

We heard about an example of a decision–implementation gap from Hatim Rahman, a professor at Northwestern University, who described a hospital project where his PhD student Jodie Koh is studying how AI is being introduced to improve access to medical diagnostics. Training the AI models requires thousands of ultrasound images, especially repeat scans of the same patients over time. But efficiency pressures in healthcare have taught clinicians to minimize the number of images for each patient to save time and money. Obtaining patient consent to use their images creates additional delays. And progress is further slowed because imaging departments that need to collaborate to implement AI tools often do not. Rahman told us they have little incentive to do so—and sometimes have a history of nasty interdepartmental conflict. Technicians also resist the AI project because they fear the resulting data will be used to monitor their



performance, increase workloads, or justify job cuts. As a result, the process of assembling the necessary images is dragging on far longer than leaders expected when the implementation began.

Investor James Cham, a partner at Bloomberg Beta, told us about an AI-native startup that tried to avoid this kind of drag. They cordoned off a small team from the rest of the company, moved them into a separate office, and gave them a clear mandate: Build end-to-end workflows that software could run without human involvement. In practice, this meant shifting the work of “management” from overseeing people to overseeing AI-based tools that performed the work themselves—tools that, for example, could review, write, and fix code based on bugs detected in production. (Matt Beane, a professor at UC Santa Barbara, and colleagues [describe](#) such work as managing “fleets” of autonomous agents.) The separation and isolation of this team at the AI startup allowed them to move faster than the rest of the organization, but it’s too soon to tell whether that speed will translate into something the broader company can adopt and sustain.

To balance the rush to move fast with the discipline to slow down where it matters, try these:



Protect the slow mode. Build speed bumps into creative and strategic work, including checkpoints, incubation periods, and reflection rituals. Perry Klebahn leads the Stanford d.school's Launchpad accelerator, where more than 100 companies have been founded in the past decade. Klebahn reports that AI makes it far easier than in the past for the Launchpad founders he coaches to generate prototypes and iterate quickly. But that very ease and speed, Klebahn says, sometimes dampens the deep commitment, determination, and ownership that founders need to develop promising ideas into viable companies. He explains, "Founders used to say 'my idea' or 'our prototype.' Now they say 'this idea.'"

This change troubles Klebahn because he believes that, when ideas come too easily, people stop fighting for them. It's a modern twist on what psychologists call the "labor leads to love" effect: The harder we work for something, the more we value it in order to justify our efforts to ourselves and others. The key is knowing which parts of the process benefit from fast execution and which require the slow thinking and struggle that enable original ideas to reach their potential.

Reward learning, not showmanship. Rebecca Stern, Udemy's senior director of learning and leadership development, told us about the company's organization-wide AI learning events, called "U-Days." Instead of just celebrating flashy demos, Udemy splits



prizes across three categories: highest business impact, most measurable improvement, and strongest peer feedback. The goal isn't to reward who moves fastest and is most dazzling, but to reward people with ideas and who give feedback that can propel deeper change.

Run the “AI residue” test. A hallmark of every management and technology fad is that vendors, founders, and internal teams litter their language with fashionable jargon to disguise shallow thinking, weak ideas, and untested solutions. Unfortunately, the term “AI” is used all too often in this way. We suggest that, after you hear a pitch, slow down, cross out every use of the term AI (and related jargon monoxide such as “cutting-edge algorithms,” “leveraging big data,” and “scalable solutions”). Read it again. If what's left is empty and hollow, congratulations—you can “pass” on the pitch and save your time and money.

Tension 5: Top-Down vs. Peer-Driven Change

Another tension we saw was how much AI transformation should be led from the top vs. driven by peers. Top-down policies, principles, and leadership are essential for fueling company-wide use and learning of AI tools (or any other systemic change). [Data from Worklytics](#) shows that teams were *twice as likely* to adopt AI tools when their managers used them first. Yet, when there's too much top-down pressure for any change, it can spark resistance



and performative compliance: employees checking boxes rather than changing how they work.

On the other hand, too much reliance on bottom-up efforts and energy, and you get fragmentation and fatigue—pockets of experimentation that never scale. One CTO compared it to having hundreds of little speed boats racing in different directions, slowing down the company because their efforts aren't coordinated. The best leaders strike a balance, setting direction from the top while encouraging and rewarding bottom-up efforts.

To balance top-down direction with enthusiasm and commitment from people throughout your organization, try these:

Create a rhythm of change. At a Fortune 20 retailer, the CEO keeps AI as a standing topic in his monthly meeting with hundreds of VPs. A cross-functional steering committee meets regularly to align adoption, governance, and use cases across finance, HR, tech, and operations. And departmental staff meetings end with the same ritual: a brief “AI moment,” where leaders take turns sharing what they’ve tried, what’s working, and what’s not. The result: AI initiatives flow through every level of the organization and have become part of the organization’s operating rhythm, not a side project.



Plan for most experiments to fail. One VP at a Fortune 500 organization told us that roughly 80% of AI projects fall short of their initial productivity goals. So, they don't redesign or eliminate jobs until *after* they have convincing evidence that AI will increase efficiency and reliability.

We suggest that as leaders are driving top-down change, they treat a high failure rate as a normal, expected, and even welcome part of doing something new. Bake in review cycles to dissect and capture what worked, what didn't, and why from employees at lower levels—so that failure fuels learning rather than frustration and finger-pointing.

Measure what matters. As part of top-down driven change efforts, we're seeing too many companies measure AI *activity*—clicks, logins, or prompt counts—instead of *impact*. For example, one Fortune 100 company even ranked employees by AI-tool usage and sent weekly leaderboards. Another tied “AI activity” to performance reviews.

Nan Guo, senior VP of engineering at Zendesk, shared how instead of tracking superficial “usage” indicators like logins or number of prompts, her team built a balanced scorecard of six engineering productivity metrics: five are operational (cycle time, code review cycle time, merge frequency, change failure rate, and number of deploys) and one is an engagement metric that



captures how engineers feel about their tools. In our experience, focusing on impact rather than activity gives people clearer guidance on what leaders value top down, enabling teams from the bottom up to adopt AI in ways that actually improve their work.

Formalize peer networks. At Uber, enterprise gen AI strategy lead Manjari Agochiya told us that they launched an open call for AI use cases. It surfaced about 150 ideas, but more importantly, it revealed 53 early AI champions across functions. These were the employees who were already experimenting, sharing lessons, and building prototypes. They became Uber's first real network of AI champions, forming the backbone of its internal peer-learning community.

In other organizations that we've studied, these kinds of champion networks often boost engagement, strengthen cross-functional morale, and open up new opportunities for visibility and advancement—especially for the non-technical employees who often emerge as early champions. At one organization, for example, a sales leader who began experimenting with AI in his day-to-day work quickly became the go-to internal expert and was later tapped to help lead the company's broader AI adoption efforts.



Leading Well in the Murky World of AI

AI transformation is a balancing act, and leaders and organizations need to wrestle with competing forces. Smart leaders don't mindlessly swing to one extreme or the other. They stay curious and treat tensions as design features and constraints—something to be managed and exploited—not as flaws or pesky problems they can and should eliminate.

Finally, predicting exactly how AI will and should shape organizations is a fool's errand. The leaders who succeed won't be ones who hold strong, inflexible, and unchanging beliefs. Research and practical lessons about implementing AI in the workplace are tentative, messy, in constant flux, and (all too often) contradictory. Any conclusion about the best path forward is, at best, provisional.

The challenge for leaders reminds us of an old (and likely apocryphal) story about a Harvard Medical School dean's warning to incoming medical students on the first day of school: "Fifty percent of what we teach you will turn out to be wrong. The problem is we don't know which 50%." Similarly, when it comes to AI, wise and winning leaders will be those who are humble enough, and build organizations that are flexible enough, so they and their teams can keep learning and updating as the ground shifts under their feet.



RH

Rebecca Hinds is the Head of the Work AI Institute at Glean. She previously founded the Work Innovation Lab at Asana and is the author of the forthcoming book, *Your Best Meeting Ever: 7 Principles for Designing Meetings That Get Things Done* (Simon & Schuster, 2026).



Robert I. Sutton is an organizational psychologist and Professor Emeritus of Management Science and Engineering at Stanford University. He has written eight books, including (with Huggy Rao) *The Friction Project: How Smart Leaders Make the Right Things Easier and the Wrong Things Harder* (St. Martin's Press, 2024).

